

This is the latest in my series of [screencasts](#) demonstrating how to use the [tidymodels](#) packages, from just starting out to tuning more complex models with many hyperparameters. Today's screencast walks through how to train and evaluate a random forest model, with this week's [#TidyTuesday dataset](#) on Mario Kart world records. 🍄

Here is the code I used in the video, for those who prefer reading instead of or in addition to video.

Explore data

Our modeling goal is to predict whether a [Mario Kart world record](#) was achieved using a shortcut or not.

```
library(tidyverse)
```

```
records <- read_csv("https://raw.githubusercontent.com/rfordatascience/tidytuesday/master/data/2021/2021-05-25/records.csv")
```

How are the world records distributed over time, for the various tracks?

```
records %>%  
  ggplot(aes(date, time, color = track)) +  
  geom_point(alpha = 0.5, show.legend = FALSE) +  
  facet_grid(rows = vars(type), cols = vars(shortcut), scales =  
    "free_y")
```



The record times decreased at first but then have been more stable. The record times are different for the different tracks, and for three lap vs. one lap times.

Build a model

Let's start our modeling by setting up our "data budget."

```
library(tidymodels)

set.seed(123)
mario_split <- records %>%
  select(shortcut, track, type, date, time) %>%
  mutate_if(is.character, factor) %>%
  initial_split(strata = shortcut)

mario_train <- training(mario_split)
mario_test <- testing(mario_split)

set.seed(234)
mario_folds <- bootstraps(mario_train, strata = shortcut)
mario_folds

## # Bootstrap sampling using stratification
## # A tibble: 25 x 2
##   splits          id
##
## 1 Bootstrap01
## 2 Bootstrap02
## 3 Bootstrap03
## 4 Bootstrap04
## 5 Bootstrap05
## 6 Bootstrap06
## 7 Bootstrap07
## 8 Bootstrap08
## 9 Bootstrap09
## 10 Bootstrap10
## # ... with 15 more rows
```

For this analysis, I am tuning a decision tree model. Tree-based models are very low-maintenance when it comes to data preprocessing, but single decision trees can be pretty easy to overfit.

```
tree_spec <- decision_tree(
  cost_complexity = tune(),
  tree_depth = tune()
) %>%
  set_engine("rpart") %>%
  set_mode("classification")

tree_grid <- grid_regular(cost_complexity(), tree_depth(), levels = 7)

mario_wf <- workflow() %>%
  add_model(tree_spec) %>%
  add_formula(shortcut ~ .)

mario_wf

## == Workflow ==
##
## Preprocessor: Formula
## Model: decision_tree()
##
## — Preprocessor —
##
## shortcut ~ .
##
## — Model —
##
## Decision Tree Model Specification (classification)
##
## Main Arguments:
##   cost_complexity = tune()
##   tree_depth = tune()
##
## Computational engine: rpart
```

Let's tune the tree parameters to find the best decision tree for this Mario Kart data set.

```
doParallel::registerDoParallel()

tree_res <- tune_grid(
  mario_wf,
  resamples = mario_folds,
  grid = tree_grid,
  control = control_grid(save_pred = TRUE)
)

tree_res

## # Tuning results
## # Bootstrap sampling using stratification
```

```
## # A tibble: 25 x 5
##   splits          id          .metrics          .notes
##   .predictions
##
## 1
```

All done! We tried all the possible combinations of tree parameters for each resample.

Choose and evaluate final model

Now we can explore our tuning results.

```
collect_metrics(tree_res)

## # A tibble: 98 x 8
##   cost_complexity tree_depth .metric .estimator mean      n
std_err .config
##
## 1 0.0000000001          1 accuracy binary    0.637    25
0.00371 Preproces...
## 2 0.0000000001          1 roc_auc  binary    0.637    25 0.0109
Preproces...
## 3 0.00000000316        1 accuracy binary    0.637    25
0.00371 Preproces...
## 4 0.00000000316        1 roc_auc  binary    0.637    25 0.0109
Preproces...
## 5 0.00000001          1 accuracy binary    0.637    25
0.00371 Preproces...
## 6 0.00000001          1 roc_auc  binary    0.637    25 0.0109
Preproces...
## 7 0.00000316          1 accuracy binary    0.637    25
0.00371 Preproces...
## 8 0.00000316          1 roc_auc  binary    0.637    25 0.0109
Preproces...
## 9 0.0001              1 accuracy binary    0.637    25
0.00371 Preproces...
## 10 0.0001             1 roc_auc  binary    0.637    25 0.0109
Preproces...
## # ... with 88 more rows

show_best(tree_res, metric = "accuracy")

## # A tibble: 5 x 8
##   cost_complexity tree_depth .metric .estimator mean      n std_err
##   .config
## 1 0.00316          8 accuracy binary    0.738    25 0.00248
```

Preprocess...

```
## 2 0.0000000001      8 accuracy binary    0.736    25 0.00249
```

Preprocess...

```
## 3 0.00000000316    8 accuracy binary    0.736    25 0.00249
```

Preprocess...

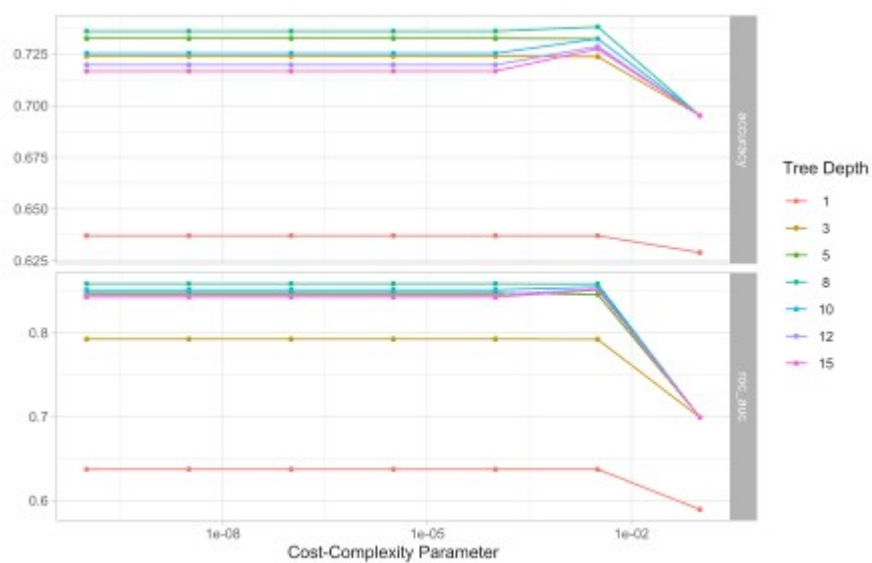
```
## 4 0.0000001        8 accuracy binary    0.736    25 0.00249
```

Preprocess...

```
## 5 0.00000316       8 accuracy binary    0.736    25 0.00249
```


Preprocess...

```
autoplot(tree_res)
```



Looks like a tree depth of 8 is best. How do the ROC curves look for the resampled training set?

```
collect_predictions(tree_res) %>%
  group_by(id) %>%
  roc_curve(shortcut, .pred_No) %>%
  autoplot() +
  theme(legend.position = "none")
```



Let's choose the tree parameters we want to use, finalize our (tuneable) workflow with this choice, and then **fit** one last time to the training data and **evaluate** on the testing data. This is the first time we have used the test set.

```
choose_tree <- select_best(tree_res, metric = "accuracy")

final_res <- mario_wf %>%
  finalize_workflow(choose_tree) %>%
  last_fit(mario_split)

collect_metrics(final_res)

## # A tibble: 2 x 4
##   .metric .estimator .estimate .config
##
## 1 accuracy binary      0.721 Preprocessor1_Model1
## 2 roc_auc  binary      0.847 Preprocessor1_Model1
```

One of the objects contained in `final_res` is a fitted workflow that we can save for future use or deployment (perhaps via `readr::write_rds()`) and use for prediction on new data.

```
final_fitted <- final_res$.workflow[[1]]
predict(final_fitted, mario_test[10:12, ])

## # A tibble: 3 x 1
##   .pred_class
##
## 1 No
## 2 No
## 3 Yes
```

We can use this fitted workflow to explore model explainability as well. Decision trees are pretty explainable already, but we might, for example, want to see a partial dependence plot for the shortcut probability and time. I like using the [DALEX](#) package for tasks like this, because it is very fully featured and has [good support for tidymodels](#). To use DALEX with tidymodels, first you create an explainer and then you use that explainer for the task you want, like computing a PDP or Shapley explanations.

Let's start by creating our "explainer."

```
library(DALEXtra)

mario_explainer <- explain_tidymodels(
  final_fitted,
  data = dplyr::select(mario_train, -shortcut),
  y = as.integer(mario_train$shortcut),
  verbose = FALSE
)
```


Then let's compute a partial dependence profile for time, grouped by `type`, which is three laps vs. one lap.

```
pdp_time <- model_profile(
  mario_explainer,
  variables = "time",
  N = NULL,
  groups = "type"
)
```

You can use the default plotting from DALEX by calling `plot(pdp_time)`, but if you like to customize your plots, you can access the underlying data via `pdp_time$agr_profiles` and `pdp_time$cp_profiles`.

```
as_tibble(pdp_time$agr_profiles) %>%
  mutate(`_label_` = str_remove(`_label_`, "workflow_")) %>%
  ggplot(aes(`_x_`, `_yhat_`, color = `_label_`)) +
  geom_line(size = 1.2, alpha = 0.8) +
  labs(
    x = "Time to complete track",
    y = "Predicted probability of shortcut",
```

```
color = NULL,  
title = "Partial dependence plot for Mario Kart world records",  
subtitle = "Predictions from a decision tree model"  
)
```



The shapes that we see here reflect how the decision tree model makes decisions along the time variable.